

Lecture 6: Data Acquisition

James Sears*

AFRE 891 SS 24

Michigan State University

*Parts of these slides are adapted from [“Data Science for Economists”](#) by Grant McDermott and [“Advanced Data Analytics”](#) by Nick Hagerty.

Table of Contents

1. **Prologue**
2. **Intro to Web Scraping**
3. **Scraping Static Websites**

Prologue

Data Acquisition

In order to wrangle, clean, or visualize data, we first need... data.



Where Do Data Come from?

There is a whole spectrum, from DIY to plug-and-chug.

1. Pre-cleaned datasets posted on secondary repositories

- General and journal replication repositories
 - e.g. [Harvard Dataverse](#)
- Community-based repositories
 - e.g. [Kaggle](#)
- Public GitHub repositories
 - e.g. [Johns Hopkins COVID-19 caseloads, deaths, and vaccination data](#)

Where Do Data Come from?

There is a whole spectrum, from DIY to plug-and-chug.

2. Open data libraries

- US Gov't: [Data.gov](https://data.gov)
- Institutions, non-profits, and thinktanks
 - e.g. [World Bank Open Data](#), [Pew Research](#), [NBER Public Use Data](#), [Economic Policy Institute](#), [AEA Data Sources](#)
- Tech and other sites
 - e.g. [data.world](#), [Our World in Data](#), [Stanford Large Network Datasets](#)
- [Map of 2600+ Open Data Portals Worldwide](#)

Where Do Data Come from?

There is a whole spectrum, from DIY to plug-and-chug.

3. Websites of primary data providers

- Government agencies; some private companies and NGOs; scientific researchers.
 - i.e. [EPA + USGS Water Quality Portal](#)

4. Programmatic data access through Application Program Interfaces (APIs)

- We'll talk more about these next lecture

Where Do Data Come from?

There is a whole spectrum, from DIY to plug-and-chug.

5. Liberate previously inaccessible data

- Build relationships with people in government or the private sector.
- Find the right person, cold-email them and ask politely.
 - i.e. my master's thesis!
- File a Freedom of Information Act (FOIA) request (a last resort; very aggressive move).

Where Do Data Come from?

6. Compile data yourself

- Assemble systematic information from many disparate sources.
- E.g. historical archives, websites, PDF reports.

7. Collect your own primary data

- Run surveys or experiments.

Where to Look for Data?

There is **no "one-stop shop"**. Where to look entirely depends on your topic.

For economics research:

- **Search the literature:** Find papers related to your topic and check the Data section.
 - Good for learning the "standard" sources for common things (e.g., weather data).
- **Find your own data** that the literature hasn't used yet.
 - It's hard to find a novel use for an already widely-used dataset.
 - Cross-disciplinary arbitrage
- **Combine data in new ways:** most new projects will require joining data from 2+ sources.
 - E.g. state-level policy changes + household-level outcome data.

Where to Look for Data?

There is **no "one-stop shop"**. Where to look entirely depends on your topic.

A few useful starting points:

- ["How to Find Data: Tips for Finding Data"](#) (Davidson College Library).
- ["Data Sets for Quantitative Research"](#) (University of Missouri Libraries).
- [Google Dataset Search](#), [AWS Data Exchange](#)

The majority of data sources described above have the data easily accessible once found.

If the data aren't already machine readable, then we can take advantage of...

Intro to Web Scraping

Web Scraping

Web scraping is the process of **extracting semi-structured web data** and **converting into a structured dataset**

- Useful when information is already online but not available in a handy format.

The screenshot displays a Yelp business profile for Meijer, a supermarket and grocery store in East Lansing, MI. The page includes a search bar at the top with the query 'grocery store' and the location 'East Lansing, MI'. The business details section on the left lists the phone number (517) 885-9053, a link to the website, a map and directions to 7157 E Saginaw St, and a link to write a review. Below this is a sponsored advertisement for Amazon.com. The 'Hours' section indicates regular hours from 9:00 am to 9:00 pm, Monday through Sunday. A 'Recommended videos' section is also present. The main content area on the right features a large image of the store's interior, the business name 'Meijer', and a description of its offerings as a family-owned, one-stop shop. It also lists 'Extra Phones', 'Payment method' (Master Card, Visa, Amex, Apple Pay, Debit, Discover, Company Card, Cash, Check), and 'AKA' (Meijer Bakery, Meijer Deli).

Yelp The Real Yellow Pages

grocery store East Lansing, MI Find

(517) 885-9053

Visit Website

Map & Directions
7157 E Saginaw St
East Lansing, MI 48823

Write a Review

Is this your business?
Customize this page. Claim This Business

Hours
Regular Hours
Mon - Sun: 9:00 am - 9:00 pm

Recommended videos
Powered by YouTube

Here are six tips for how to be a great Next Stay

Meijer
Supermarkets & Super Stores, Grocery Stores
☆☆☆☆☆ Be the first to review!
OPEN NOW Today: 9:00 am - 9:00 pm
89 Years in Business

Shop Groceries Online Today - Free Grocery Pickup
Sponsored - https://www.amazon.com/ +
Enjoy free curbside or in-store pickup at your local Whole Foods Market. Prime members can shop online, save time and get free pickup later.
Visit Website

More Info

General Info
Meijer is your family-owned, one-stop shop in East Lansing, MI that's been offering our neighbors great food, great brands, and great value since 1934. Stop in for the freshest produce delivered daily from local growers, custom-cut quality meats, seafood delivered 6 days a week, bread baked fresh daily, plus low prices across 40+ departments, from premium pet brands to the latest fashions for the entire family. Plus, save even more with weekly specials and mPerks.

Extra Phones
Phone: (517) 885-9010
Phone: (517) 885-9047
Phone: (517) 339-3611
Phone: (517) 885-9000
Fax: (517) 885-9010

Payment method
Master Card, Visa, Amex, Apple Pay, Debit, Discover, Company Card, Cash, Check

AKA
Meijer Bakery
Meijer Deli

Structure of Webpages

Webpages are largely made out of **two types of files** that we have to parse:

HTML

- HyperText Markup Language is a **markup language**, Like Markdown
- It specifies the **structure** of a webpage.

CSS

- Cascading Style Sheets is a language for **formatting the appearance of a webpage**
 - CSS **properties** specify **how** to format: what font, what color, how wide
 - CSS **selectors** specify **what** to format: which structural elements get what rule

How Websites Render Content

It's also worth realising that there are **two ways** that web content gets rendered in a browser:

1. **Server-side (back-end)**
2. **Client-side (front-end)**

You can read [here](#) for more details (including example scripts), but for our purposes the essential features are as follows...

Server-Side Content Rendering

The scripts that build a **server-side** website aren't run on our computer

- Rather, the script that "builds" the site is **run on a host server**
 - All information is .hi-blue[directly embedded] in the webpage's HTML
- e.g. Wikipedia tables are already populated with all the info we see in our browser

No. ♦	Peak ♦	Range (or island) ♦	Location ♦	Coordinates ^[1] ♦	Prominence (m) ♦	Height (m) ♦	Col (m) ♦
1.	Mount Everest	Himalayas	 China  Nepal	 27°59'17.6500°N 86°55'30.0652°E	8,848.86	8,848.86	0
2.	Aconcagua	Andes	 Argentina	 32°39'11"S 70°0'42"W	6,960.8	6,960.8	0
3.	Denali / Mount McKinley*	Alaska Range	 United States	 63°4'10"N 151°0'26"W	6,155	6,191	47
4.	Mount Kilimanjaro*	Eastern Rift mountains	 Tanzania	 3°4'0"S 37°21'33"E	5,885	5,895	10
5.	Pico Cristóbal Colón	Sierra Nevada de Santa Marta	 Colombia	 10°50'18"N 73°41'12"W	5,509	5,700	191

Server-Side Content Rendering

The scripts that build a **server-side** website aren't run on our computer

- Rather, the script that "builds" the site is **run on a host server**
 - All information is .hi-blue[directly embedded] in the webpage's HTML
- e.g. Wikipedia tables are already populated with all the info we see in our browser

- **Webscraping process:** finding the correct selectors (CSS or Xpath), iterating through (dynamic) webpages (e.g. "Next page" and "Show more" tabs)
- **Key concepts:** CSS, Xpath, HTML

Client-Side Content Rendering

The scripts that build a **client-side** website aren't run on our computer

- Website contains an empty template of HTML and CSS
 - May contain a "skeleton" table without any values
- When we visit the page URL, our browser sends a **request** to the host server
- If everything is okay (e.g. our request is valid), then the server sends a **response** script, which our browser executes and uses to populate the HTML template with the specific information that we want.

Client-Side Content Rendering

- If everything is okay (e.g. our request is valid), then the server sends a **response** script, which our browser executes and uses to populate the HTML template with the specific information that we want.

[illegible][illegible]

Web Scraping

Over the next two lectures we'll cover the main differences between the two approaches and general workflows.

However, I want to forewarn that web scraping typically involves a fair bit of **detective work**, iterating and adjusting steps

- According to the type of data you want
- To match the specifics of a given website

In short, web scraping involves **as much art as it does science.**

The good news, though: **if you can see it, you can scrape it.**

Caveat: Ethical and Legal

The last line brings up an important consideration: just because you **can** scrape it, doesn't mean you **should**.

Legality

- In short, it is **currently legal** to scrape data from the web using automated tools as long as the data are **publicly available** (hiQ Labs vs. LinkedIn)
 - We'll chat more later on about what this means for "hidden" APIs
- May get blocked due to

Ethicality

- Need to consider the impact your scraper will have on the host server
 - Easy to write a function that can overwhelm a website's host with rapid requests
 - We'll return to the "be nice" mantra later on

Scraping Static Websites

Static Scraping: Preliminaries

Today we'll be using [SelectorGadget](#), which is a Chrome extension that makes it easy to discover CSS selectors.

- Install the extension directly [here](#).

Please note that SelectorGadget is only available for **Chrome**. If you prefer using **Firefox**, then you can try [ScrapeMate](#).

Static Scraping: Preliminaries

The primary R package that we'll be using today is rvest, a simple web scraping library inspired by Python's Beautiful Soup, but with extra tidyverse functionality. **rvest** is designed to work with webpages that are built server-side and thus requires knowledge of the relevant CSS selectors...

Which means that now is probably a good time for us to cover what these are in more detail. First, let's load some packages

```
pacman::p_load(lubridate, rvest, tidyverse)
```


CSS Selectors and SelectorGadget

CSS **selectors** are the **"what"** of the display rules. They identify which rules should be applied to which elements.

- E.g. Text elements that are selected as **".h1"** (i.e. top line headers) are usually larger and displayed more prominently than text elements selected as **".h2"** (i.e. sub-headers).

The key point is that if you can **identify the CSS selector(s)** of the content you want, then you can **isolate it from the rest** of the webpage content that you don't want.

CSS Selectors and SelectorGadget

This where SelectorGadget comes in. We'll work through an extended example (with a twist!) below, but I highly recommend looking over this [quick vignette](#) soon.

Here are two helpful links if you're interested in reading more about [CSS](#) (i.e Cascading Style Sheets) and [SelectorGadget](#)

Application 1: Wikipedia

Let's say you watched the U.S. Olympic Marathon Trials earlier this month and now want to scrape the wikipedia page on [marathon world record progression](#)

- Women's record: 3:40:22 in 1926 to 2:11:53 last year!

First, open up this page in your browser. Take a look at its structure: What type of objects does it contain? How many tables does it have? Do these tables all share the same columns? What row- and columns-spans? Etc.

Application 1: Wikipedia

Contents

hide

(Top)

Criteria for record eligibility

Women's world record

Unofficial record attempts

History

Men

Women

Gallery of world record holders

See also

Notes

References

Sources

External links

Marathon world record progression

ArticleTalk

From Wikipedia, the free encyclopedia

This list is a chronological progression of record times for the marathon. *World records in the marathon* are ratified by *World Athletics*, the international governing body for the sport of athletics.

Kenyan athlete *Kevin Kiptum* set a men's world record time of 2:00:35 on October 8, 2023, at the *2023 Chicago Marathon*.^[1]

Ethiopian athlete *Tigst Assefa* broke the women's world record for a mixed-gender race with a time of 2:11:53 on September 24, 2023, at the *2023 Berlin Marathon*.^[2] In addition to the standard women's marathon world record, World Athletics also recognizes a second world record for women in the "Women Only" category, meaning that the marathon was run on a course without any male athletes in competition. The current "Women Only" record of 2:17:01 was set by *Mary Keitany* on April 23, 2017, at the London Marathon in the elite women's race.^[3]

Criteria for record eligibility

See also: *List of world records in athletics § Criteria*

For a performance to be ratified as a world record by World Athletics, the marathon course on which the performance occurred must be 42.195 km (26.219 mi) long,^[4] measured in a defined manner using the *calibrated bicycle method*^[5] (the distance in kilometers being the official distance, the distance in miles is an approximation) and meet other criteria that rule out artificially fast times produced on courses aided by downhill slope or tailwind.^[6] The criteria include:

- "The start and finish points of a course, measured along a theoretical straight line between them, shall not be further apart than 50% of the race distance."^[4]
- "The decrease in elevation between the start and finish shall not exceed an average of one in a thousand, i.e. 1m per km."^[4]

In recognizing Kenyan *Geoffrey Mutai*'s mark of 2:03:02 at the *2011 Boston Marathon* as (at the time) "the fastest Marathon ever run", the IAAF said: "Due to the elevation drop and point-to-point measurements of the Boston course, performances [on that course] are not eligible for World record consideration."^[6]


Road racing events like the marathon were specifically excepted from World Athletics rule 260.18(d) that rejected from consideration those track and field performances set in mixed competition.^[6]

The *Association of Road Racing Statisticians*, an independent organization that compiles data from *road running* events, also maintains an alternate marathon world best progression but with standards they consider to be more stringent.^{[10][11]}


Women's world record

The IAAF Congress at 2011 World Championships in Athletics passed a motion changing the record eligibility criteria effective October 6, 2007, so that women's world records must be set in all-women competitions.^[12] The result of the change was that Radcliffe's 2:17:42 performance at the 2005 London Marathon would supplant her own existing women's mark as the "world record"; the earlier performance was to be referred to as a "world best".^[12]

The decision was met with strong protest in Britain. In November 2011, an IAAF council member announced that Radcliffe's original mark would be allowed to stand, with the eventual decision that both marks would be



Kevin Kiptum during his world record run at the 2023 Chicago marathon with 2:00:35



Tigst Assefa during her women's

Time	Name	Nationality	Date	Event/Place	Source	Notes
5:40:xx	Marie-Louise Ledru	 France	September 29, 1918	Tour de Paris Marathon	ARRS ^[10]	
3:40:22	Violet Piercy	 United Kingdom	October 3, 1926	London ^[nb 7]	IAAF ^[53]	The ARRS indicates that Piercy's 3:40:22 was set on August 2, 1926, during a time trial on a course that was only 35.4 km ^[10]
3:37:07	Merry Lepper	 United States	December 16, 1963 ^[nb 8]	Culver City, United States	IAAF ^[53]	Disputed (short course) ^[95]
3:27:45	Dale Greig	 United Kingdom	May 23, 1964	Ryde	IAAF ^[53] ARRS ^[10]	
3:19:33	Mildred Sampson	 New Zealand	July 21, 1964 ^[nb 8]	Auckland, New Zealand	IAAF ^[53]	Disputed by ARRS as a time trial ^{[nb 9][96]}
3:14:23	Maureen Wilton	 Canada	May 6, 1967	Toronto, Canada	IAAF ^[53] ARRS ^[10]	The ARRS notes Wilton's extended time as 3:14:22.8 ^[10]
3:07:27.2	Anni Pedersen	 West Germany	September 16, 1967	Waldniel, West Germany	IAAF ^[53] ARRS ^[10]	The ARRS notes Pedersen's extended time as 3:07:26.2 ^[10]
3:02:53	Caroline Walker	 United States	February 28, 1970	Seaside, OR	IAAF ^[53] ARRS ^[10]	
3:01:42	Elizabeth Bonner	 United States	May 9, 1971	Philadelphia, United States	IAAF ^[53] ARRS ^[10]	
2:55:22	Elizabeth Bonner	 United States	September 19, 1971	New York City Marathon	IAAF ^[53] ARRS ^[10]	
2:49:40	Cheryl Bridges	 United States	December 5, 1971	Culver City, United States	IAAF ^[53] ARRS ^[10]	
2:46:36	Michiko Gorman	 United States	December 2, 1973	Culver City, United States	IAAF ^[53] ARRS ^[10]	The ARRS notes Gorman's extended time as 2:46:37 ^[10]
2:46:24	Chantal Langlacé	 France	October 27, 1974	Neuf-Brisach, France	IAAF ^[53] ARRS ^[10]	
2:43:54.5	Jacqueline Hansen	 United States	December 1, 1974	Culver City, United States	IAAF ^[53] ARRS ^[10]	The ARRS notes Hansen's extended time as 2:43:54.6 ^[10]

Application 1: Wikipedia

Once you've familiarised yourself with the structure, read the whole page into R using the `rvest::read_html()` function.

```
mthn = read_html("https://en.wikipedia.org/wiki/Marathon_world_record_pro  
mthn
```

```
## {html_document}  
## <html class="client-nojs vector-feature-language-in-header-enabled vector-fe  
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=UTF-  
## [2] <body class="skin-vector skin-vector-search-vue mediawiki ltr sitedir-lt
```

Application 1: Wikipedia

As you can see, this is an XML document that contains **everything** needed to render the Wikipedia page.¹

It's kind of like viewing someone's entire LaTeX document (preamble, syntax, etc.) when all we want are the data from some tables in their paper.

¹ XML stands for Extensible Markup Language and is one of the primary languages used for encoding and formatting web pages.

First Table: Women's Records

Let's start by scraping our first table from the page, which documents the women's record progression.


The first thing we need to do is identify the table's unique CSS selector using SelectorGadget.

Note: this *will* require trial and error - and a lot of clicking.

First Table: Women's Records

Start by activating SelectorGadget and clicking on a chart element.

2:01:09	Eliud Kipchoge	 Kenya	September 25, 2022	Berlin Marathon	World Athletics ^[90]	
2:00:35	Kelvin Kiptum	 Kenya	October 8, 2023	Chicago Marathon	World Athletics ^[91]	First man to break 2:01:00 in a record-eligible marathon.



Listed by World Athletics as a world best prior to official acceptance^[53]
Ratified by World Athletics as a world best (since January 1, 2003) or world record (since January 1, 2004)^[53]
Recognized by the Association of Road Racing Statisticians (ARRS)^[10]

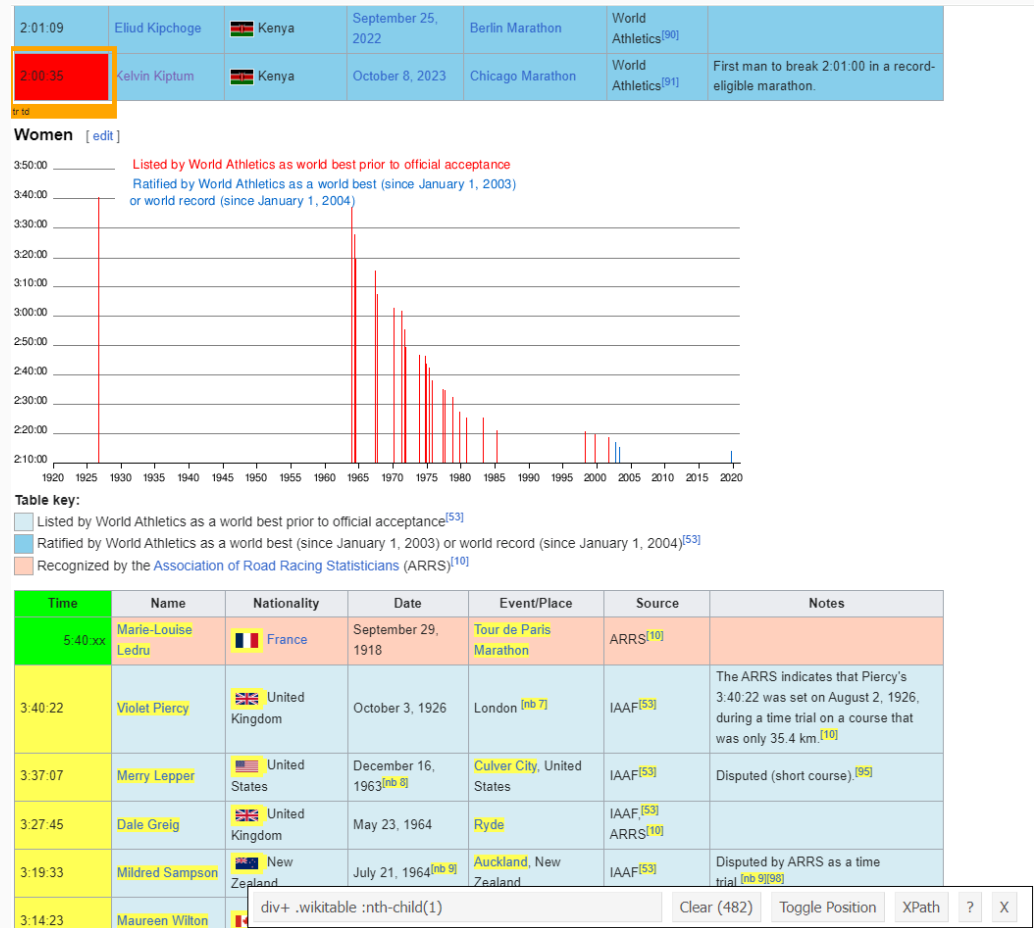
Time	Name	Nationality	Date	Event/Place	Source	Notes
5:40:xx	Marie-Louise Ledru	 France	September 29, 1918	Tour de Paris Marathon	ARRS ^[10]	
3:40:22	Violet Piercy	 United Kingdom	October 3, 1926	London ^[nb 7]	IAAF ^[53]	The ARRS indicates that Piercy's 3:40:22 was set on August 2, 1926, during a time trial on a course that was only 35.4 km. ^[10]
3:37:07	Merry Lepper	 United States	December 16, 1963 ^[nb 8]	Culver City, United States	IAAF ^[53]	Disputed (short course). ^[95]
3:27:45	Dale Greig	 United Kingdom	May 23, 1964	Ryde	IAAF ^[53] ARRS ^[10]	
3:19:33	Mildred Sampson	 New Zealand	July 21, 1964 ^[nb 9]	Auckland, New Zealand	IAAF ^[53]	Disputed by ARRS as a time trial. ^{[nb 9][96]}
3:14:23	Maureen Wilton	 Canada				

Clicking on another chart element expands the selection (but too much!)



First Table: Women's Records

Click the elements you *don't* want until they turn red:



First Table: Women's Records

Working through this iterative process yields

```
"div+ .wikitable :nth-child(1)".
```

We can use this unique CSS selector to isolate the women's record table content from the rest of the HTML document.

- **Extract the table content** with `html_element()`
- Parse the **HTML table into an R data frame** with `html_table()`

First Table: Women's Records

- **Extract the table content** with `html_element()`
- Parse the **HTML table into an R data frame** with `html_table()`

```
women <- mthn %>%  
  html_element("div+ .wikitable :nth-child(1)") %>% ## select table element  
  html_table() ## convert to data frame
```

women

```
## # A tibble: 43 × 7
```

	Time	Name	Nationality	Date	Event/Place	Source No
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
## 1	5:40:xx	Marie-Louise Ledru	France	Septe...	Tour de Pari...	ARRS[... ']
## 2	3:40:22	Violet Piercy	United Kingdom	Octob...	London [nb 7]	IAAF[... ']
## 3	3:37:07	Merry Lepper	United States	Decem...	Culver City,...	IAAF[... ']
## 4	3:27:45	Dale Greig	United Kingdom	May 2...	Ryde	IAAF,...
## 5	3:19:33	Mildred Sampson	New Zealand	July ...	Auckland, Ne...	IAAF[... ']
## 6	3:14:23	Maureen Wilton	Canada	May 6...	Toronto, Can...	IAAF,...
## 7	3:07:27.2	Anni Pede-Erdkamp	West Germany	Septe...	Waldniel, We...	IAAF,...
## 8	3:02:53	Caroline Walker	United States	Febru...	Seaside, OR	IAAF,...

First Table: Women's Records

Great, it worked! Now convert the date string to a format that R actually understands.²

```
women %>%  
  mutate(Date = mdy(Date)) %>% ## convert string to date format  
  head(4)
```

```
## # A tibble: 4 × 7
```

##	Time	Name	Nationality	Date	Event/Place	Source	No
##	<chr>	<chr>	<chr>	<date>	<chr>	<chr>	<
## 1	5:40:xx	Marie-Louise Ledru	France	1918-09-29	Tour de Pari...	ARRS[...	'
## 2	3:40:22	Violet Piercy	United Kingd...	1926-10-03	London [nb 7]	IAAF[...	'
## 3	3:37:07	Merry Lepper	United States	NA	Culver City,...	IAAF[...	'
## 4	3:27:45	Dale Greig	United Kingd...	1964-05-23	Ryde	IAAF,...	'

² Note: If column name had spaces or capital letters, we could use the

`janitor::clean_names()` convenience function to clean them. (Q: How else could we have done this?)

First Table: Women's Records

Alright that mostly worked, but there are a few hyperlink references in the dates leading to NAs.

- This is a case where using a regular expression is convenient: match the pattern "[nb X]" at the **end of the strings**

```
women <- women %>%  
  mutate(across(where(is.character), # mutate across all character string.  
    ~str_replace(.x, "\\[nb [0-9]+\\]$", "")), # remove [nb 0  
    Date = mdy(Date)) # convert string to date format
```

women

```
## # A tibble: 43 × 7
```

##	Time	Name	Nationality	Date	Event/Place	Source	No
##	<chr>	<chr>	<chr>	<date>	<chr>	<chr>	<
##	1 5:40:xx	Marie-Louise Led...	France	1918-09-29	"Tour de Par...	ARRS[...	'
##	2 3:40:22	Violet Piercy	United Kin...	1926-10-03	"London "	IAAF[...	'
##	3 3:37:07	Merry Lepper	United Sta...	1963-12-16	"Culver City...	IAAF[...	'
##	4 3:27:45	Dale Greig	United Kin...	1964-05-23	"Ryde"	IAAF,...	'

Table 2: Men's Records

We could stop here and plot the women's records, but while we're here let's grab the men's records as well.

Challenge: take a couple minutes to use SelectorGadget and find the CSS selector for the men's record table.

- Don't peek at next slide until you give it a shot!

```
men <- mthn %>%  
  html_element("") %>% ## add the selector!  
  html_table() %>%  
  mutate(across(where(is.character), # mutate across all character string.  
                ~str_replace(.x, "\\[nb [0-9]+\\]$", "")), # remove [nb 0  
    Date = mdy(Date))
```

Table 2: Men's Records

```
men ← mthn %>%  
  html_element("p+ .wikitable :nth-child(1)") %>% ## add the selector!  
  html_table() %>%  
  mutate(across(where(is.character), # mutate across all character string.  
    ~str_replace(.x, "\\[nb [0-9]+\\]$", "")), # remove [nb 0  
    Date = mdy(Date))  
tail(men)
```

```
## # A tibble: 6 × 7
```

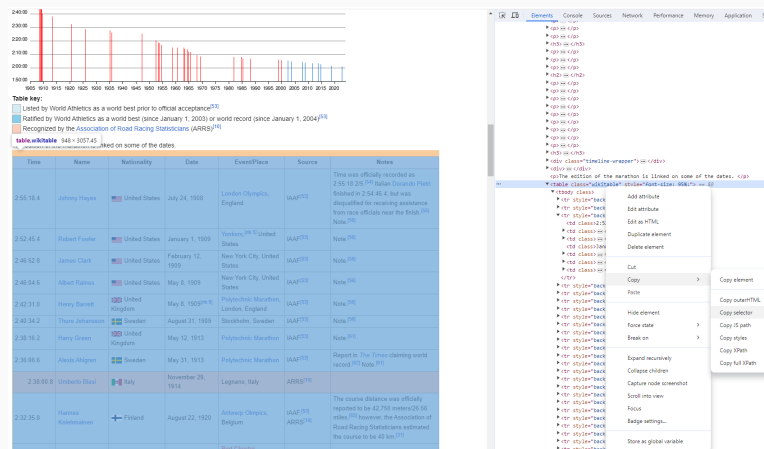
##	Time	Name	Nationality	Date	Event/Place	Source	No
##	<chr>	<chr>	<chr>	<date>	<chr>	<chr>	<
## 1	2:03:38	Patrick Makau	Kenya	2011-09-25	Berlin Marathon	IAAF,[82...	'
## 2	2:03:23	Wilson Kipsang	Kenya	2013-09-29	Berlin Marathon	IAAF[85]...	'
## 3	2:02:57	Dennis Kimetto	Kenya	2014-09-28	Berlin Marathon	IAAF[87]...	'
## 4	2:01:39	Eliud Kipchoge	Kenya	2018-09-16	Berlin Marathon	IAAF[89]	'
## 5	2:01:09	Eliud Kipchoge	Kenya	2022-09-25	Berlin Marathon	World At...	'
## 6	2:00:35	Kelvin Kiptum	Kenya	2023-10-08	Chicago Marathon	World At...	'

Browser Inspection Tools

SelectorGadget is a great tool, but sometimes it **takes more work than necessary** and isn't available in all browsers.

Alternate approach: use the inspect web element

- Chrome: Right-click > "Inspect"
(`Ctrl + Shift + I`)
- Scroll over source elements until the table of interest is highlighted
- Use the selector that pops up over the element
 - i.e. `"table.wikitable"`



Duplicate Selectors

If we look at the source code a bit more we can see that the selector `table.wikitable` **isn't unique!**

In cases like this, use `html_elements()` to retrieve all matching elements as a list

```
mthn_tabs ← mthn %>%  
  html_elements("table.wikitable") ## select all elements matching the se  
  
# first match is men's results  
men2 ← mthn_tabs[[1]] %>% html_table()  
  
# second match is women's results  
women2 ← mthn_tabs[[2]] %>% html_table()
```

Marathon Progression

Challenge: combine both tables into a single dataframe and plot the record progression

- Clean the `Time` variable (get rid of alphanumeric characters at the end)
- Separate the time variable into hours/minutes/seconds using `separate()`
- Use `hours()`, `minutes()`, `seconds()` from `lubridate` to create a time variable for the marathon time
- Create a plot with
 - Marathon time on the y-axis
 - Make sure to include a `scale_y_time()` layer!
 - Date on the x-axis
 - Linetype aesthetic mapping based on the group (men vs. women)
 - Use a nice theme

Answer in a few slides (no peeking until you've tried first)


Marathon Progression

```
men <- mutate(men, Group = "Men")
women <- mutate(women, Group = "Women")

mthn_prog <- rbind(men, women) %>%
  mutate(Time = str_replace_all(Time, "[:alpha:]", "")) %>%
  separate(col = "Time", into = c("Hours", "Minutes", "Seconds"), sep = ' ')
  # add time variable using lubridate functions
  mutate(mthn_time = hours(Hours) + minutes(Minutes) + seconds(Seconds))

ggplot(mthn_prog, aes(x = Date, y = mthn_time)) +
  geom_line(aes(linetype = Group)) +
  scale_y_time() +
  labs(y = "Marathon World Record",
       x = "Year") +
  theme_minimal()
```

Marathon Progression



Scraping Static Sites

Next time we'll see another case of how to get specific text elements that *aren't* in a nice table format before talking about considerations for dynamically-generated sites and APIs.

Table of Contents

1. **Prologue**
2. **Intro to Web Scraping**
3. **Scraping Static Websites**